


Real-Time Low-Latency Full-Duplex Audio Streaming Server

Author: Santhi Maddipudi



Problem we are solving

Voice AI sounds robotic because most implementations are turn-based.

You talk, it waits, then it responds.

Real humans interrupt each other mid-sentence.

Current voice bots cannot handle this.

They feel unnatural and frustrating to use.

Tech stack

FastAPI, WebSockets, Silero VAD, Deepgram, OpenAI Realtime API, ElevenLabs, Python asyncio.

Proof of work :

In progress

